

Homework 0: Alohomora!

Marco Huang

Department of Computer Science
University of Maryland, College Park
hhuang01@umd.edu

I. PHASE-1

”Shake My Boundary,” ventures into the realm of boundary detection by implementing a simplified Pb detection algorithm that examines brightness, color, and texture gradients. It also compares the results against traditional edge detection methods.

A. Filter Banks

DoG, Leung-Malik, and Gabor Filters are commonly used methods for feature extraction in image processing and computer vision. These filter banks help in detecting and representing different orientations and scales in images, which are essential for tasks such as boundary detection and texture analysis.

1) *Oriented DoG filters*: Created by convolving a simple Sobel filter and a Gaussian kernel with sigma = [5, 10] kernel size = 50 orientations = 16

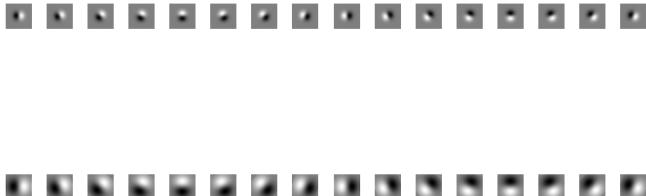


Fig. 1: Oriented Derivative of Gaussian (DoG) filters.

2) *Leung-Malik Filters*: LM Small (LMS), the filters occur at basic scales $\sigma = \{1, \sqrt{2}, 2, 2\sqrt{2}\}$. The first and second derivative filters occur at the first three scales with an elongation factor of 3, i.e., ($\sigma_x = \sigma$ and $\sigma_y = 3\sigma_x$).

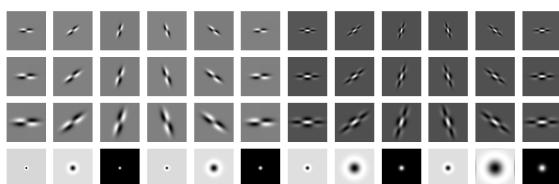


Fig. 2: Laplacian of Gaussian (LOG) filters

3) *Gabor Filters*: Gabor filter is a gaussian kernel function modulated by a sinusoidal plane wave. $\text{sigma3} = [3, 5, 7, 9, 12]$ orientations = 8 wavelength = 5 phase offset = 0 aspect ratio = 0.5

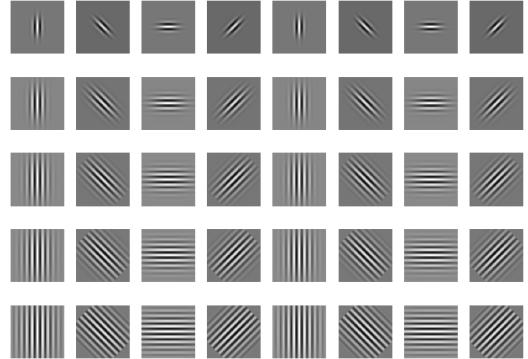


Fig. 3: Gabor Filters

B. Texton, Brightness, Color Map

An image is filtered through an N-filter array, turning each pixel into an N-dimensional vector based on filter responses. KMeans clustering then assigns each pixel a distinct texton ID, resulting in a Texton map. Brightness and Color maps are similarly derived from pixel grayscale values and RGB values, respectively.

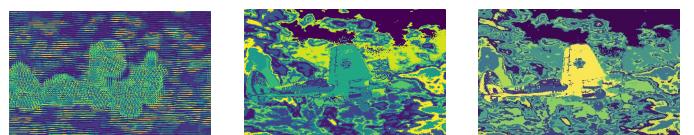


Fig. 4: Texton, brightness and color maps for image 1

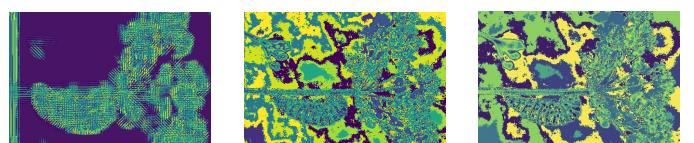


Fig. 5: Texton, brightness and color maps for image 2

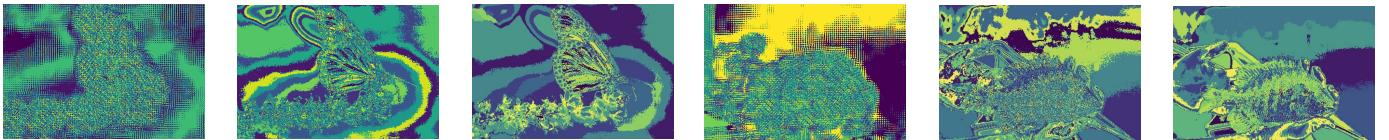


Fig. 6: Texton, brightness and color maps for image 3

Fig. 12: Texton, brightness and color maps for image 9

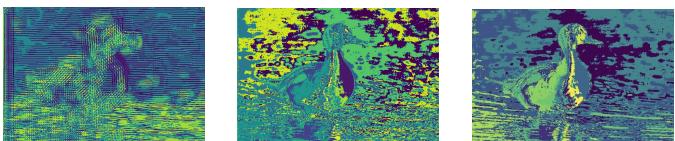


Fig. 7: Texton, brightness and color maps for image 4

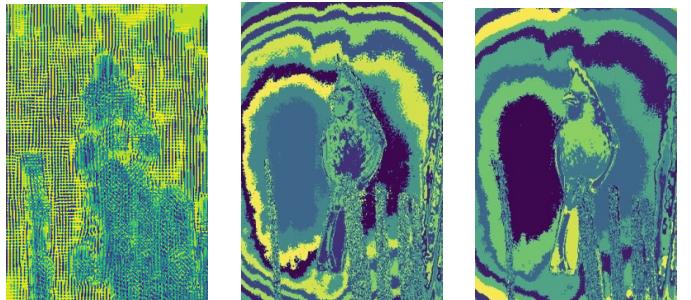


Fig. 13: Texton, brightness and color maps for image 10

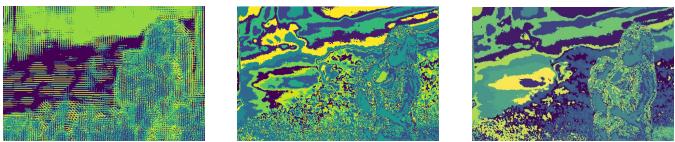


Fig. 8: Texton, brightness and color maps for image 5

C. Texture, brightness, color gradient map

To compute the texture, brightness, and color gradient map. We need to compute differences of values across different shapes and sizes. This can be achieved very efficiently by the use of Half-disc masks.

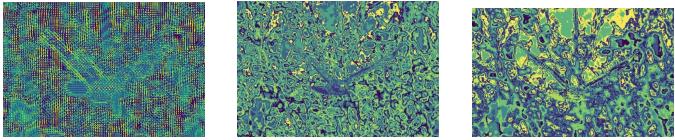


Fig. 9: Texton, brightness and color maps for image 6

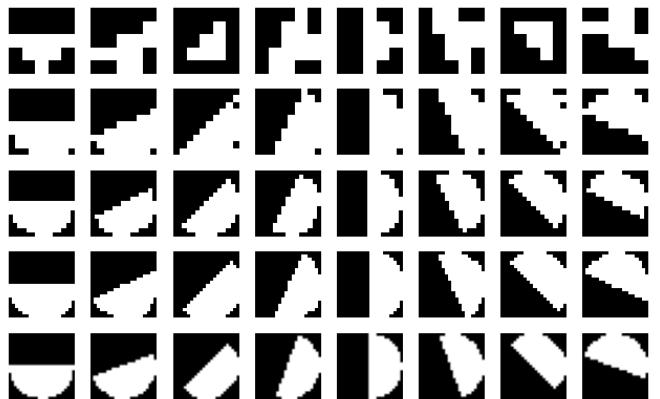


Fig. 14: Half-disc masks

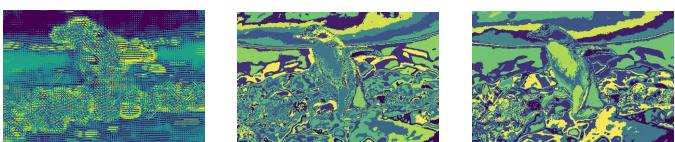


Fig. 10: Texton, brightness and color maps for image 7

Using these half-disc masks along with the Chi-square distance using a filtering operation. We compute by comparing the distributions in left/right half-disc pairs

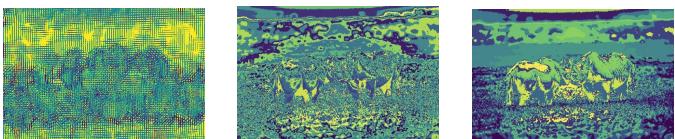


Fig. 11: Texton, brightness and color maps for image 8

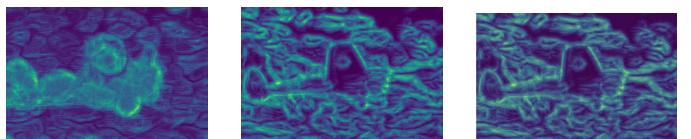


Fig. 15: Gradient maps for image 1

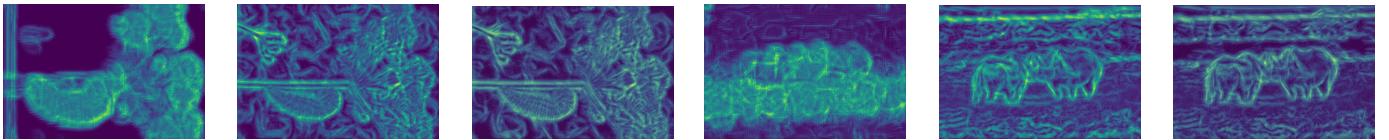


Fig. 16: Gradient maps for image 2

Fig. 22: Gradient maps for image 8

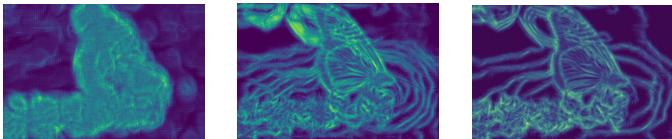


Fig. 17: Gradient maps for image 3

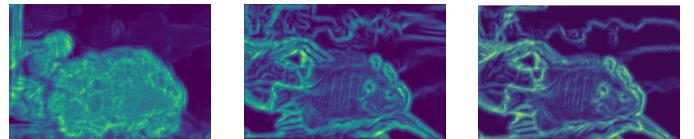


Fig. 23: Gradient maps for image 9

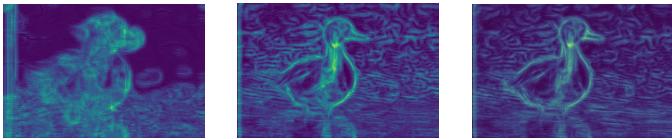


Fig. 18: Gradient maps for image 4

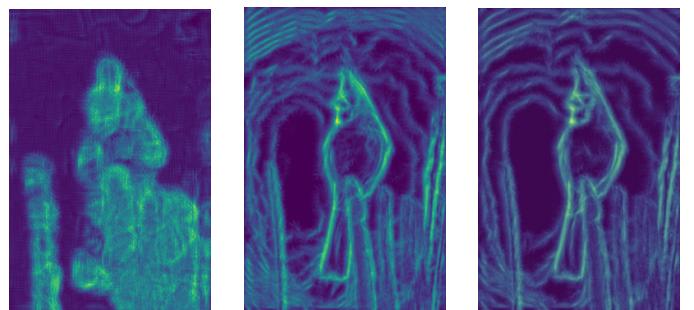


Fig. 24: Gradient maps for image 10

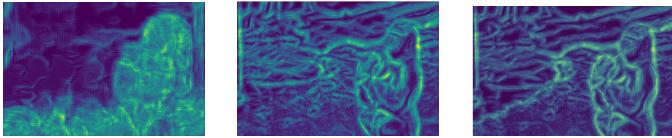


Fig. 19: Gradient maps for image 5

D. Pb-lite Output

The final step is to combine information from the features with a baseline method (based on Sobel or Canny edge detection or an average of both) using a simple equation

$$PbEdges = \frac{(T_g + B_g + C_g)}{3} \odot (w_1 * cannyPb + w_2 * sobelPb)$$

Here I choose $w_1 = 0.5$ and $w_2 = 0.5$. Compare Canny baseline, Sobel baseline and Pb-lite output.

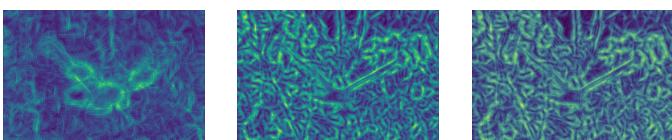


Fig. 20: Gradient maps for image 6



Fig. 25: Canny, Sobel and Pblite for image 1

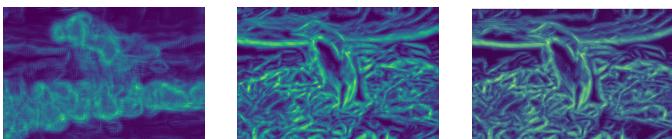


Fig. 21: Gradient maps for image 7



Fig. 26: Canny, Sobel and Pblite for image 2



Fig. 27: Canny, Sobel and Pblite for image 3



Fig. 33: Canny, Sobel and Pblite for image 9

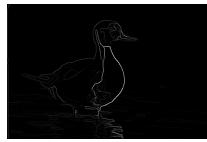


Fig. 28: Canny, Sobel and Pblite for image 4

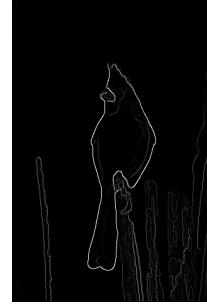


Fig. 34: Canny, Sobel and Pblite for image 10



Fig. 29: Canny, Sobel and Pblite for image 5



Fig. 30: Canny, Sobel and Pblite for image 6



Fig. 31: Canny, Sobel and Pblite for image 7



Fig. 32: Canny, Sobel and Pblite for image 8

E. Analysis

PB-Lite implementation offers a nuanced balance between sensitivity and specificity in edge detection across images with complex textures. Unlike the Canny baseline, which tends to overemphasize edges leading to numerous false positives, and the Sobel baseline, which might underplay important details by being overly conservative, PB-Lite appears to adeptly moderate edge detection. This moderation ensures that edges in highly textured areas are suppressed effectively, avoiding the pitfall of false positives without obscuring critical image features.

II. PHASE-2

“Deep Dive on Deep Learning” transitions into the application of deep learning architectures to the CIFAR-10 dataset, aiming to enhance classification accuracy. It involves training baseline neural networks, augmenting data, and experimenting with advanced architectures like ResNet, ResNeXt, and DenseNet. The report analyzes the performance and efficiency of these models, providing a comprehensive comparison of their capabilities.

A. Some detail

60000 CIFAR-10 data, 50000 training set and 10000 testing set. Using ADAM optimizer and Cross-Entropy loss function for all 5 training.

B. First neural network

The first task has no normalisation and standardization. With weight decay = 0.0001 and learning rate = 0.001. Size of the Mini Batch to use is 32. The number of parameters in this model are 12.

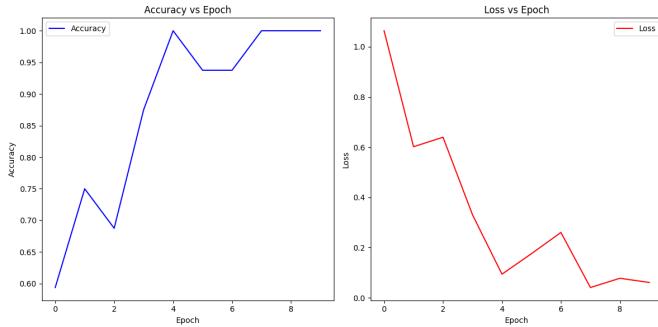


Fig. 35: Train Accuracy and Loss over Epochs for CNN basic

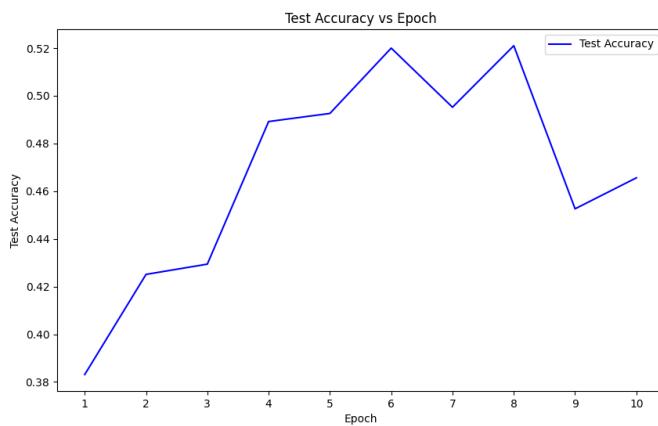


Fig. 36: Test Accuracy over Epochs for CNN basic

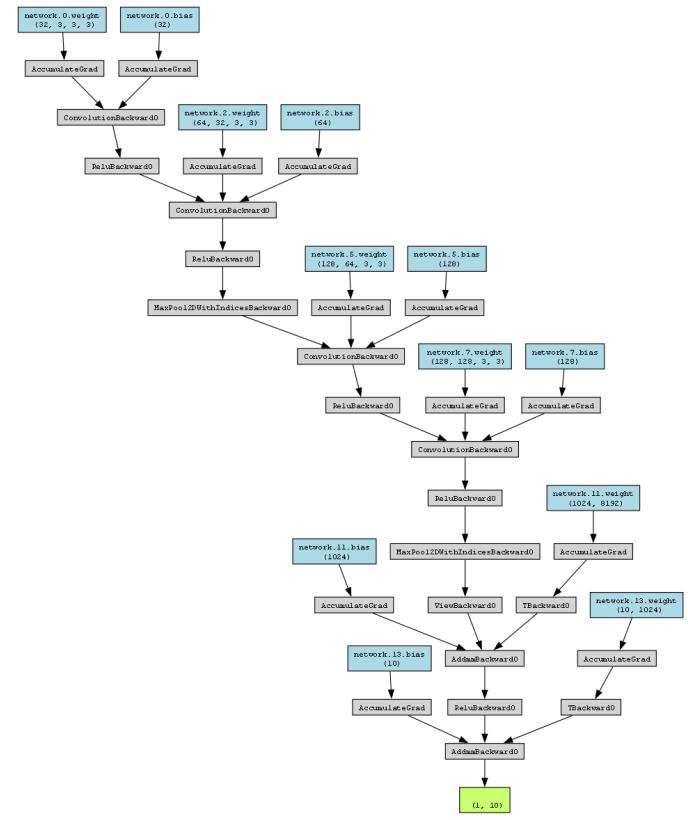


Fig. 37: CNN architecture

[3114	945	159	21	11	24	9	27	129	561]	(0)
[4	4818	5	0	0	2	0	0	1	170]	(1)
[248	704	2782	65	5	201	97	147	50	701]	(2)
[167	920	407	1412	10	630	68	128	30	1228]	(3)
[247	1489	510	135	564	214	151	373	39	1278]	(4)
[87	627	325	153	4	2837	59	183	20	705]	(5)
[86	1357	264	86	1	148	1839	37	23	1159]	(6)
[73	711	208	37	5	108	16	2986	9	847]	(7)
[209	1116	26	7	1	9	2	7	2748	875]	(8)
[12	569	4	2	0	4	0	3	7	4399]	(9)

(0) (1) (2) (3) (4) (5) (6) (7) (8) (9)

Accuracy: 54.998 %

Fig. 39: Confusion Matrix on Training data for CNN basic

[517	189	68	7	4	10	5	11	50	139]	(0)
[2	923	2	0	0	1	0	0	3	69]	(1)
[78	168	413	23	9	69	36	45	8	151]	(2)
[47	218	99	175	5	149	24	32	9	242]	(3)
[44	299	128	47	53	53	30	87	11	248]	(4)
[30	135	80	62	2	470	19	41	9	152]	(5)
[26	271	53	24	2	33	330	12	10	239]	(6)
[19	167	43	9	4	39	6	498	3	212]	(7)
[61	233	12	2	1	3	3	3	480	202]	(8)
[5	180	5	0	0	4	0	4	5	797]	(9)

Accuracy: 46.56 %

Fig. 38: Confusion Matrix on Testing data for CNN basic

C. Improving Accuracy of neural network

To improve my first neural network, I add normalisation with ((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)). Decrease learning rate to 0.005. Increase size of the Mini Batch to 64. Using tf.RandomCrop and tf.RandomHorizontalFlip. Number of parameters in this model are 12.

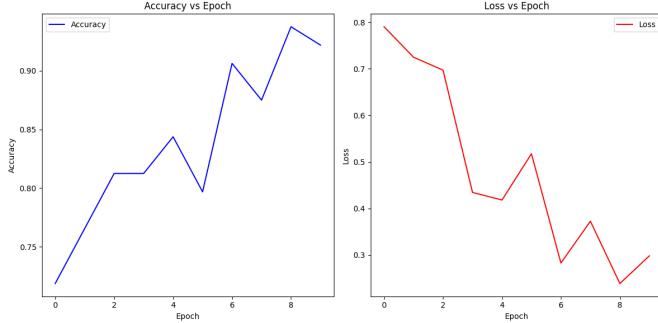


Fig. 40: Train Accuracy and Loss over Epochs for improved CNN

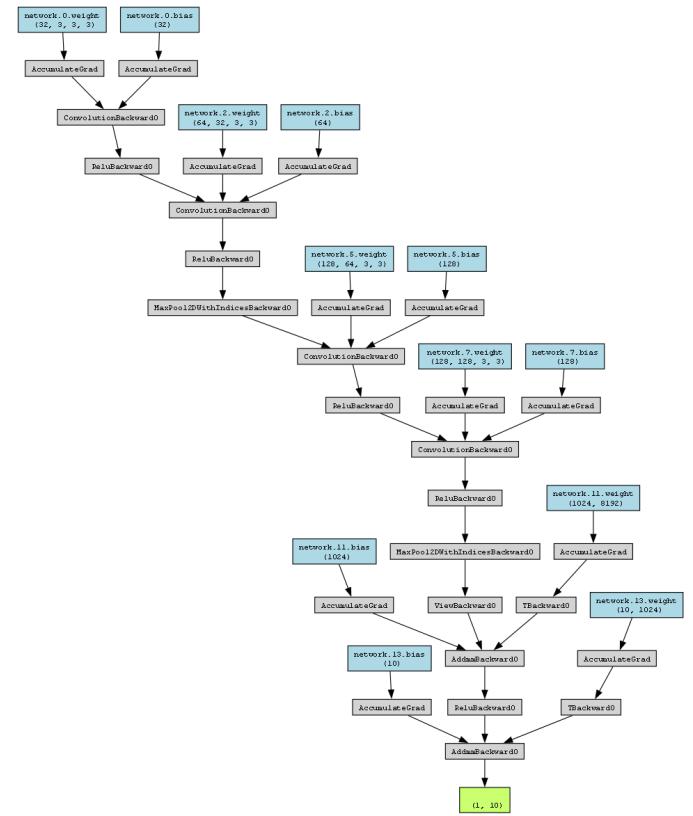


Fig. 42: Improved CNN architecture

[899	9	11	3	4	1	1	14	29	29]	(0)
[11	902	1	4	0	1	1	0	6	74]	(1)
[59	4	718	53	43	31	55	22	4	11]	(2)
[27	1	30	728	22	94	36	36	16	10]	(3)
[16	2	29	45	768	26	34	73	5	2]	(4)
[14	1	22	165	18	731	11	32	3	3]	(5)
[11	1	25	47	11	10	881	6	5	3]	(6)
[14	2	12	30	12	21	3 900	2	4]	(7)	
[50	11	2	5	2	2	2	3 908	15]	(8)	
[16	24	1	8	0	2	3	5 14 927]	(9)		
(0)	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	
Accuracy: 83.62 %										

Fig. 43: Confusion Matrix on Testing data for improved CNN

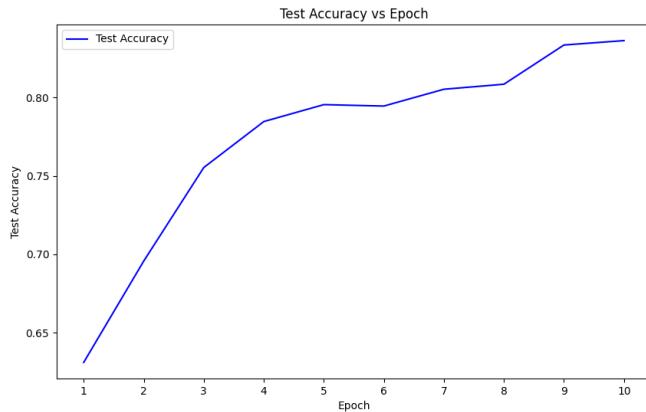


Fig. 41: Test Accuracy over Epochs for improved CNN

[4706	17	41	21	11	10	15	25	80	74]	(0)
[19	4647	3	5	0	5	5	5	29	282]	(1)
[256	12	3993	192	133	100	171	99	22	22]	(2)
[102	4	98	4042	71	379	126	122	34	22]	(3)
[66	1	84	198	4122	88	91	325	12	13]	(4)
[32	4	77	631	83	3942	42	167	12	10]	(5)
[32	8	88	141	30	32 4630	19	17	3]	(6)	
[31	0	39	94	41	69	8 4696	5	17]	(7)	
[131	39	9	13	7	3	4	5 4740	49]	(8)	
[53	54	6	14	2	12	4	16	29 4810]	(9)	
(0)	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	
Accuracy: 88.656 %										

Fig. 44: Confusion Matrix on Training data for improved CNN

D. ResNet, ResNeXt, DenseNet

Normalisation with ((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)) Learning rate is 0.005. Size of the Batch is 64

1) ResNet: The number of parameters in this model are 122

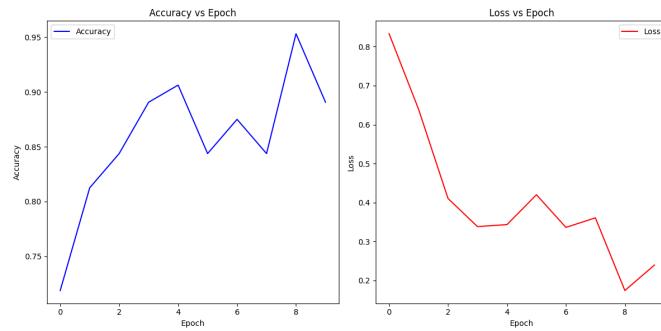


Fig. 45: Train Accuracy and Loss over Epochs for ResNet

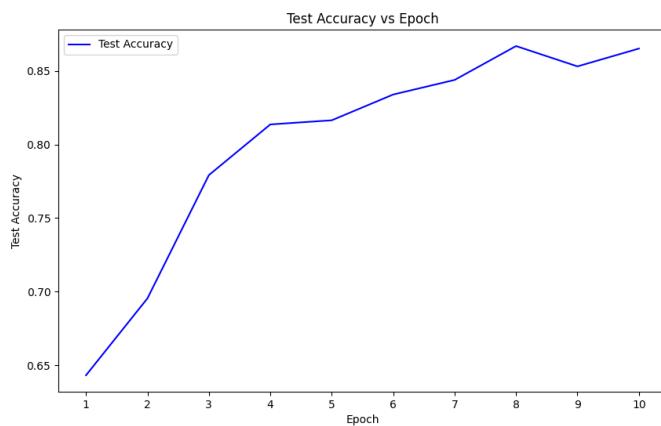


Fig. 46: Test Accuracy over Epochs for ResNet

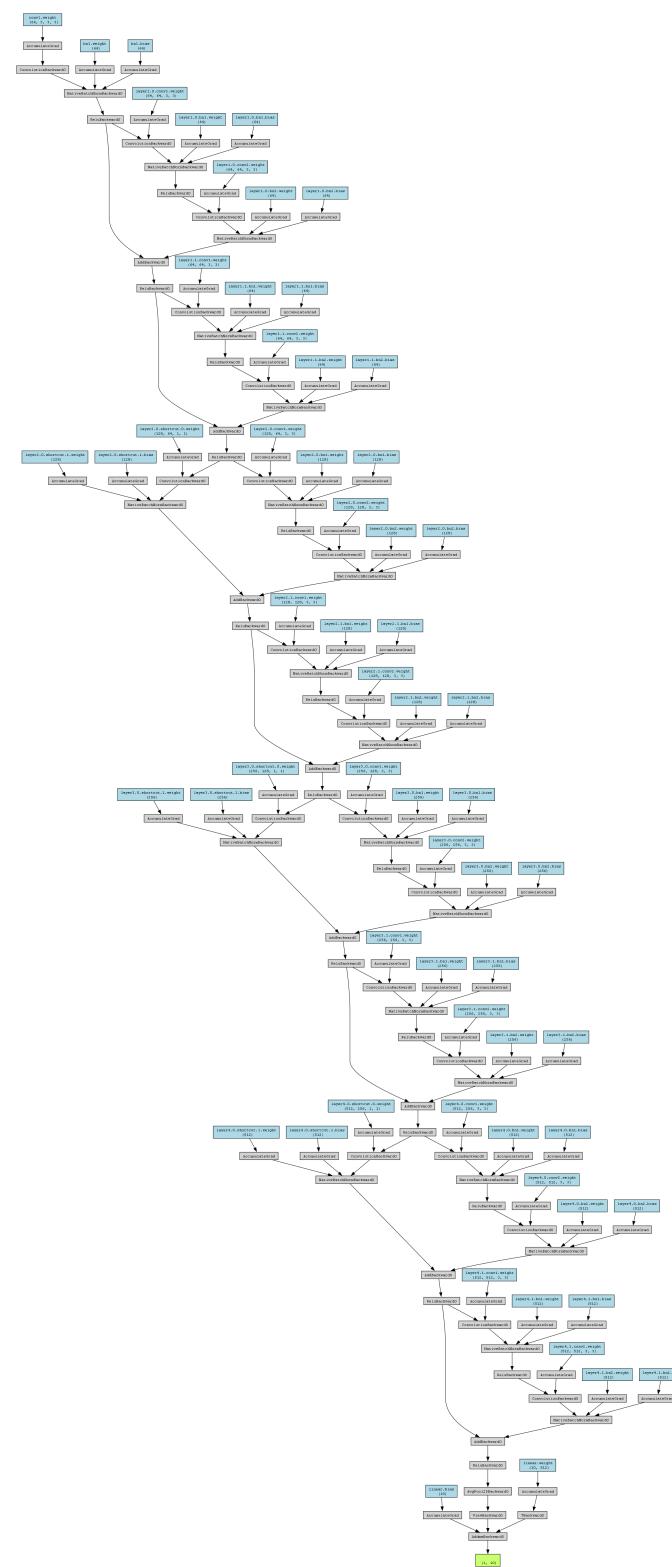


Fig. 47: ResNet architecture

[836	14	20	17	3	3	10	4	41	52]	(0)
[4	933	0	2	0	0	4	2	5	50]	(1)
[38	0	763	46	26	22	77	14	7	7]	(2)
[6	1	24	827	14	46	43	23	5	11]	(3)
[10	1	23	51	785	19	40	65	2	4]	(4)
[0	0	15	162	15	770	5	31	0	2]	(5)
[3	1	9	34	5	3	934	6	2	3]	(6)
[9	1	9	18	6	14	0	934	2	7]	(7)
[27	23	4	3	0	1	9	0	908	25]	(8)
[1	20	1	6	0	0	3	1	7	961]	(9)
(0)	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	
Accuracy: 86.51 %										

Fig. 48: Confusion Matrix on Testing data for ResNet

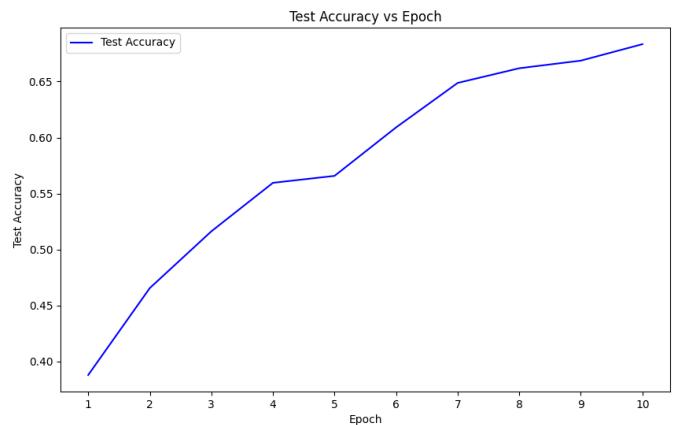


Fig. 51: Test Accuracy over Epochs for ResNeXt

[4373	53	82	88	7	5	42	12	116	222]	(0)
[4	4767	1	16	0	0	15	2	11	184]	(1)
[144	5	4083	191	68	69	330	57	19	34]	(2)
[10	3	58	4416	46	188	174	56	13	36]	(3)
[32	3	56	265	4033	66	164	337	21	23]	(4)
[7	3	59	733	39	3935	38	177	2	7]	(5)
[9	7	18	103	18	34	4788	7	10	6]	(6)
[12	2	21	118	11	49	16	4743	2	26]	(7)
[73	48	13	15	0	2	17	6	4757	69]	(8)
[10	43	4	18	0	2	18	10	12	4883]	(9)
(0)	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	
Accuracy: 89.556 %										

Fig. 49: Confusion Matrix on Training data for ResNet

2) *ResNeXt*: The number of parameters in this model are
314

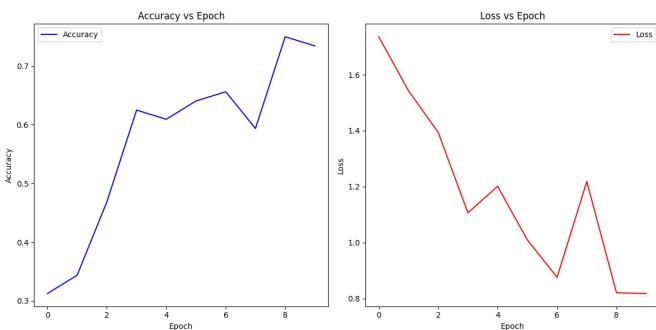


Fig. 50: Train Accuracy and Loss over Epochs for ResNeXt

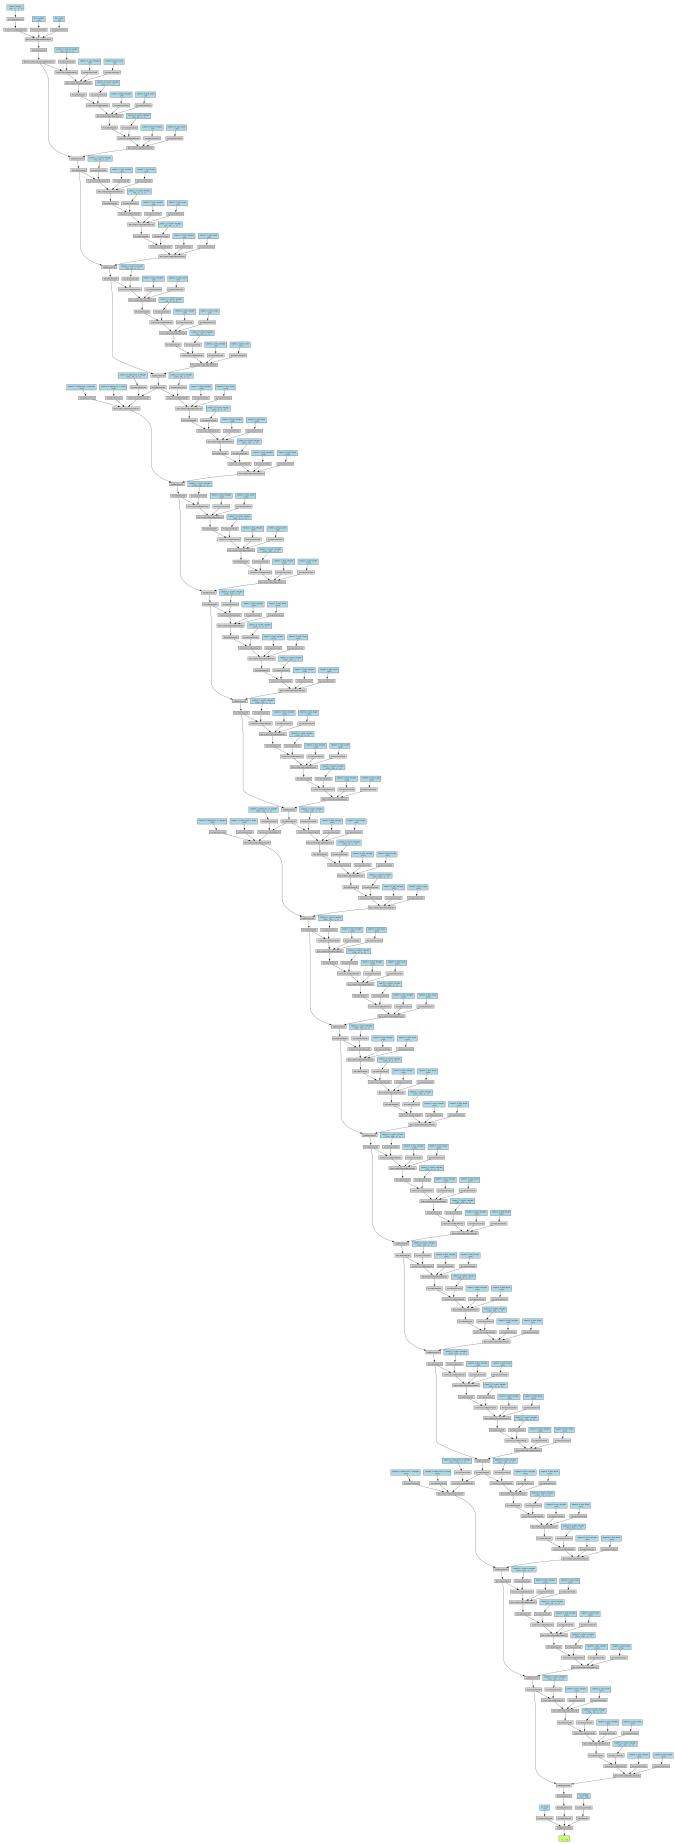


Fig. 52: ResNeXt architecture

[801	16	43	28	24	10	10	22	38	8]	(0)
[158	706	8	16	4	12	13	11	24	48]	(1)
[84	1	578	73	89	73	53	40	6	3]	(2)
[19	7	76	553	35	179	66	49	9	7]	(3)
[19	1	71	84	624	27	72	96	5	1]	(4)
[9	1	41	208	43	603	23	66	4	2]	(5)
[9	4	55	84	18	22	790	9	8	1]	(6)
[9	2	32	33	54	72	16	773	3	6]	(7)
[130	26	15	13	10	1	7	5	783	10]	(8)
[68	168	9	28	8	13	15	37	31	623]	(9)
(0)	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	
Accuracy: 68.34 %										

Fig. 53: Confusion Matrix on Testing data for ResNeXt

[4068	54	229	125	114	50	35	118	157	50]	(0)
[764	3658	32	77	14	26	77	37	138	177]	(1)
[349	14	3122	333	382	289	259	192	37	23]	(2)
[82	18	251	2994	161	870	297	258	48	21]	(3)
[138	5	264	311	3306	149	291	510	19	7]	(4)
[26	4	210	1058	170	3011	148	339	21	13]	(5)
[49	23	287	476	134	124	3856	28	15	8]	(6)
[43	6	116	196	272	272	34	4026	12	23]	(7)
[582	142	68	61	43	18	28	15	4006	37]	(8)
[346	780	38	123	35	42	79	154	115	3288]	(9)
(0)	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	
Accuracy: 70.67 %										

Fig. 54: Confusion Matrix on Training data for ResNeXt

3) DenseNet: The number of parameters in this model are 727

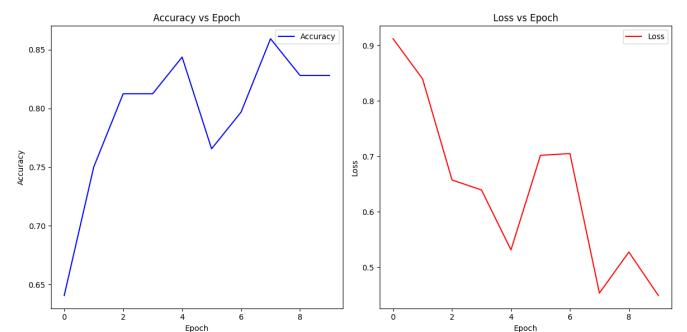


Fig. 55: Train Accuracy and Loss over Epochs for DenseNet

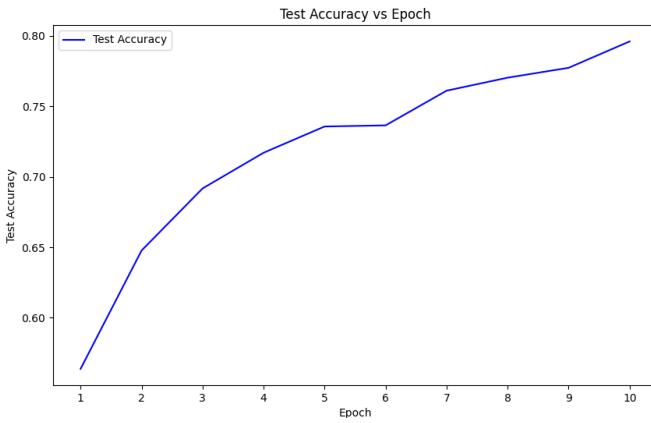


Fig. 56: Test Accuracy over Epochs for DenseNet

[4373	53	82	88	7	5	42	12	116	222]	(0)
[4	4767	1	16	0	0	15	2	11	184]	(1)
[144	5	4083	191	68	69	330	57	19	34]	(2)
[10	3	58	4416	46	188	174	56	13	36]	(3)
[32	3	56	265	4033	66	164	337	21	23]	(4)
[7	3	59	733	39	3935	38	177	2	7]	(5)
[9	7	18	103	18	34	4788	7	10	6]	(6)
[12	2	21	118	11	49	16	4743	2	26]	(7)
[73	48	13	15	0	2	17	6	4757	69]	(8)
[10	43	4	18	0	2	18	10	12	4883]	(9)
(0)	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	

Fig. 59: Confusion Matrix on Training data for DenseNet

E. Analysis

Architecture Complexity: ResNet18, ResNeXt50, and DenseNet have more complex architectures compared to the basic and improved CNN models. This complexity allows for better feature extraction and learning capabilities, leading to higher accuracies.

Data Preprocessing: The introduction of normalization and data augmentation (random crop and flip) in the improved CNN model significantly boosts its accuracy by enhancing generalization.

Learning Rate and Mini-Batch Size: The adjustment of the learning rate and mini-batch size in the improved CNN model optimizes the learning process, allowing for more stable and effective training.

Number of Parameters: Although the improved CNN model has the same number of parameters as the basic CNN model, the more advanced models (ResNet18, ResNeXt50, DenseNet) have significantly more parameters, which correlates with their ability to learn more complex patterns. However, this also means they require more computational resources.

In summary, the ResNet18 outperforms the other models due to its balance of complexity and efficiency, followed by the improved CNN model, which shows significant improvement over the basic CNN through better preprocessing and optimized training settings. DenseNet and ResNeXt50 offer different trade-offs between parameter count, accuracy, and inference time, with DenseNet achieving higher accuracy due to its efficient feature reuse mechanism.

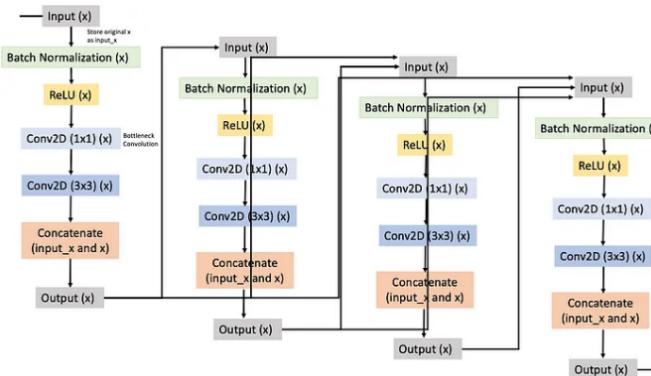


Fig. 57: DenseNet architecture, source: [1].

[841	11	20	17	9	3	6	7	43	43]	(0)
[18	874	5	4	2	5	10	1	21	60]	(1)
[49	0	688	56	71	46	48	15	15	12]	(2)
[27	1	42	690	44	109	39	23	14	11]	(3)
[8	2	42	53	761	30	30	60	11	3]	(4)
[9	1	28	206	27	667	15	34	10	3]	(5)
[9	0	40	57	24	18	838	4	9	1]	(6)
[12	0	12	35	21	49	6	843	5	17]	(7)
[52	13	7	13	0	1	3	2	885	24]	(8)
[19	61	3	16	0	3	2	5	19	872]	(9)
(0)	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	
Accuracy:	79.59	%								

Fig. 58: Confusion Matrix on Testing data for DenseNet

Model	No of Parameters	Train Accuracy	Test Accuracy	Run-time
CNN	12	54.998	46.56	0.0149
Improved CNN	12	88.656	83.62	0.0104
ResNet18	122	89.556	86.51	0.0103
ResNeXt50	314	70.670	68.34	0.0079
DenseNet	727	83.470	79.59	0.0159

Fig. 60: Analyze Table

REFERENCES

- [1] C. Martins, “Densenet architecture explained with code examples,” <https://cdanielaaam.medium.com/densenet-architecture-explained-with-code-examples-0d86d7936f83>, 2023, accessed: 2024-02-14.